

Name: Siya Dadheech

Assignment - 5

Kubernetes Installation

1. Create a Kubernetes cluster using minikube

Step 1 Launch an EC2 Instance

Step 2 Install Docker

- `sudo apt install -y docker.io`
- `sudo usermod -aG docker $USER`

Step 3 Install Kubectl

- `curl -LO "https://dl.k8s.io/release/$(curl -s https://dl.k8s.io/release/stable.txt)/bin/linux/amd64/kubectl"`
- `chmod +x kubectl`
- `sudo mv kubectl /usr/local/bin/`
- `kubectl version --client`

Step 4 Install Minikube

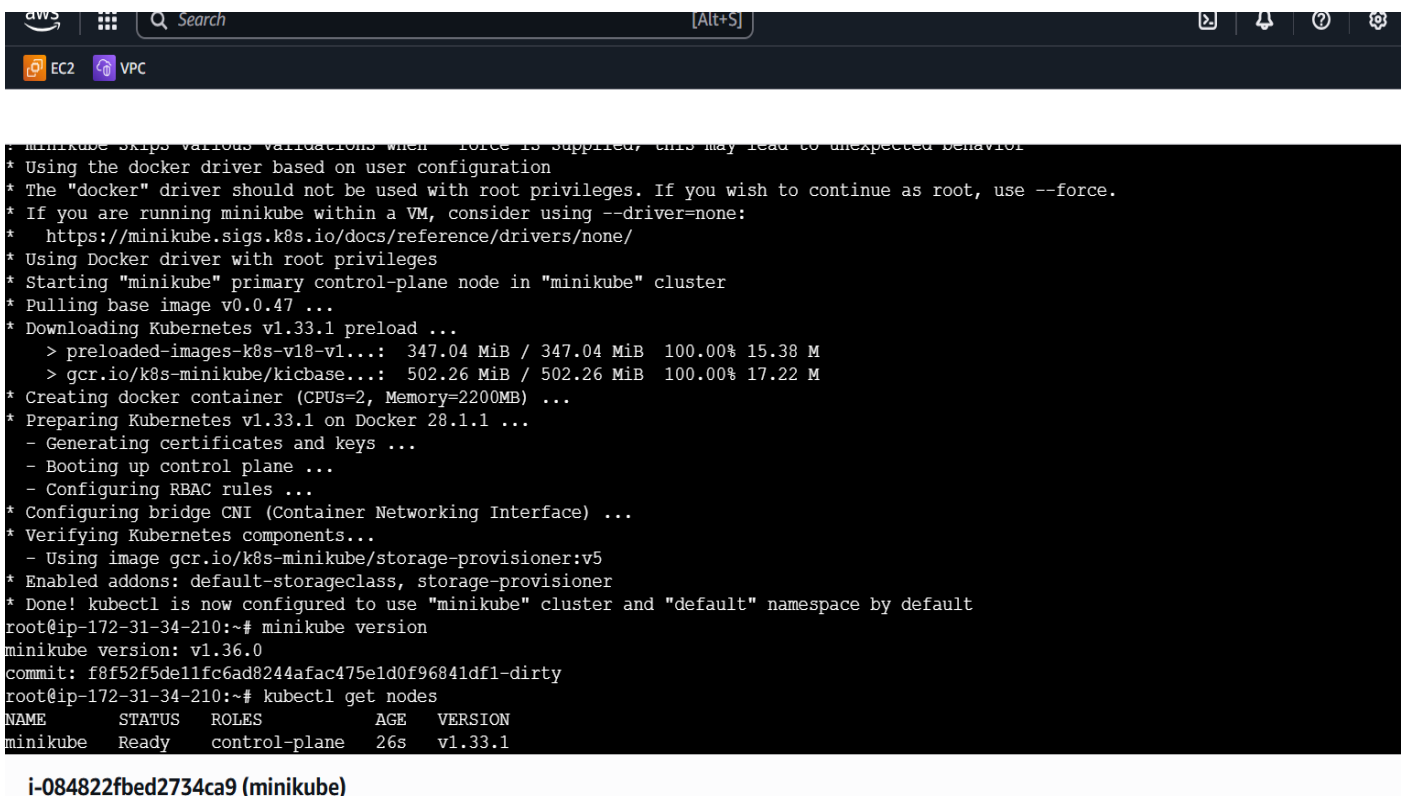
- `curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64`
- `sudo install minikube-linux-amd64 /usr/local/bin/minikube`
- `minikube version`

Step 5 To start the minikube we need driver.

- `minikube start --driver=docker`

Step 6 Verify Cluster

- `kubectl get nodes`



```
aws [Alt+S]
EC2 VPC

* minikube skips various validations when --force is supplied; this may lead to unexpected behavior
* Using the docker driver based on user configuration
* The "docker" driver should not be used with root privileges. If you wish to continue as root, use --force.
* If you are running minikube within a VM, consider using --driver=none:
*   https://minikube.sigs.k8s.io/docs/reference/drivers/none/
* Using Docker driver with root privileges
* Starting "minikube" primary control-plane node in "minikube" cluster
* Pulling base image v0.0.47 ...
* Downloading Kubernetes v1.33.1 preload ...
  > preloaded-images-k8s-v18-v1...: 347.04 MiB / 347.04 MiB 100.00% 15.38 M
  > gcr.io/k8s-minikube/kicbase...: 502.26 MiB / 502.26 MiB 100.00% 17.22 M
* Creating docker container (CPUs=2, Memory=2200MB) ...
* Preparing Kubernetes v1.33.1 on Docker 28.1.1 ...
  - Generating certificates and keys ...
  - Booting up control plane ...
  - Configuring RBAC rules ...
* Configuring bridge CNI (Container Networking Interface) ...
* Verifying Kubernetes components...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Enabled addons: default-storageclass, storage-provisioner
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
root@ip-172-31-34-210:~# minikube version
minikube version: v1.36.0
commit: f8f52f5de11fc6ad8244afac475e1d0f96841df1-dirty
root@ip-172-31-34-210:~# kubectl get nodes
NAME        STATUS    ROLES    AGE   VERSION
minikube    Ready     control-plane  26s   v1.33.1

i-084822fbed2734ca9 (minikube)
```

```

git commit.
root@ip-172-31-34-210:~# sudo usermod -aG docker $USER
root@ip-172-31-34-210:~# curl -LO "https://dl.k8s.io/release/${curl -s https://dl.k8s.io/release/stable.txt}/bin/linux/amd64/kubectl"
curl: Failed to extract a sensible file name from the URL to use for storage
curl: (3) URL using bad/illegal format or missing URL
root@ip-172-31-34-210:~# curl -LO "https://dl.k8s.io/release/${curl -L -s https://dl.k8s.io/release/stable.txt}/bin/linux/amd64/kubectl"
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100  138    100  138    0    0   453      0  --:--:-- --:--:-- --:--:--   453
100 57.3M  100 57.3M    0    0 24.7M      0  0:00:02 0:00:02 --:--:-- 30.1M
root@ip-172-31-34-210:~# chmod +x kubectl
root@ip-172-31-34-210:~# sudo mv kubectl /usr/local/bin
root@ip-172-31-34-210:~# kubectl version --client
Client Version: v1.33.2
Kustomize Version: v5.6.0
root@ip-172-31-34-210:~# ls
minikube-linux-amd64  snap
root@ip-172-31-34-210:~# curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100 126M  100 126M    0    0 11.6M      0  0:00:10 0:00:10 --:--:-- 15.0M
root@ip-172-31-34-210:~# sudo install minikube-linux-amd64 /usr/local/bin/minikube
root@ip-172-31-34-210:~# minikube version
minikube version: v1.36.0
commit: f8f52f5de11fc6ad8244afac475e1d0f96841df1-dirty
root@ip-172-31-34-210:~# sudo swapoff -a
sudo sed -i 's/swap/d' /etc/fstab

```

i-084822fbbed2734ca9 (minikube)

2. Create a Kubernetes cluster using kubeadm
 - Step 1: I Created one master and 2 worker node.
 - Step 2: both on (worker and master node)

Step 1: Disable Swap on All Nodes

swapoff -a

sed -i 's/ swap / s/^(.*)\$/#\1/g' /etc/fstab

Step 2: Enable IPv4 Packet Forwarding

cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf

net.ipv4.ip_forward = 1

EOF

Step 3: Verify IPv4 Packet Forwarding

sysctl net.ipv4.ip_forward

Step 4: Install containerd

Add Docker's official GPG key:

sudo apt-get update

sudo apt-get install ca-certificates curl

sudo install -m 0755 -d /etc/apt/keyrings

sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc

sudo chmod a+r /etc/apt/keyrings/docker.asc

Add the repository to Apt sources:

```
echo \
```

```
"deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc]  
https://download.docker.com/linux/ubuntu \
```

```
$(. /etc/os-release && echo "$VERSION_CODENAME") stable" | \
```

```
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

```
sudo apt-get update && sudo apt-get install containerd.io && systemctl enable --now containerd
```

Step 5: Install CNI Plugin

```
wget https://github.com/containernetworking/plugins/releases/download/v1.4.0/cni-plugins-linux-amd64-v1.4.0.tgz
```

```
mkdir -p /opt/cni/bin
```

```
tar Cxvf /opt/cni/bin cni-plugins-linux-amd64-v1.4.0.tgz
```

Step 6: Forward IPv4 and Configure iptables

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
```

```
overlay
```

```
br_netfilter
```

```
EOF
```

```
sudo modprobe overlay
```

```
sudo modprobe br_netfilter
```

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
sudo sysctl --system
```

```
sysctl net.bridge.bridge-nf-call-iptables net.bridge.bridge-nf-call-ip6tables net.ipv4.ip_forward
```

```
modprobe br_netfilter
```

```
sysctl -p /etc/sysctl.conf
```

Step 7: Modify containerd Configuration for systemd Support

```
Vim /etc/containerd/config.toml
```

```
disabled_plugins = []
```

```
imports = []
```

```
oom_score = 0
```

```
plugin_dir = ""
```

```
required_plugins = []  
root = "/var/lib/containerd"  
state = "/run/containerd"  
version = 2
```

```
[cgroup]  
  path = ""
```

```
[debug]  
  address = ""  
  format = ""  
  gid = 0  
  level = ""  
  uid = 0
```

```
[grpc]  
  address = "/run/containerd/containerd.sock"  
  gid = 0  
  max_recv_message_size = 16777216  
  max_send_message_size = 16777216  
  tcp_address = ""  
  tcp_tls_cert = ""  
  tcp_tls_key = ""  
  uid = 0
```

```
[metrics]  
  address = ""  
  grpc_histogram = false
```

```
[plugins]
```

```
[plugins."io.containerd.gc.v1.scheduler"]  
  deletion_threshold = 0  
  mutation_threshold = 100  
  pause_threshold = 0.02  
  schedule_delay = "0s"  
  startup_delay = "100ms"
```

[plugins."io.containerd.grpc.v1.cri"]

disable_apparmor = false
disable_cgroup = false
disable_hugetlb_controller = true
disable_proc_mount = false
disable_tcp_service = true
enable_selinux = false
enable_tls_streaming = false
ignore_image_defined_volumes = false
max_concurrent_downloads = 3
max_container_log_line_size = 16384
netns_mounts_under_state_dir = false
restrict_oom_score_adj = false
sandbox_image = "k8s.gcr.io/pause:3.5"
selinux_category_range = 1024
stats_collect_period = 10
stream_idle_timeout = "4h0m0s"
stream_server_address = "127.0.0.1"
stream_server_port = "0"
systemd_cgroup = false
tolerate_missing_hugetlb_controller = true
unset_seccomp_profile = ""

[plugins."io.containerd.grpc.v1.cri".cni]

bin_dir = "/opt/cni/bin"
conf_dir = "/etc/cni/net.d"
conf_template = ""
max_conf_num = 1

[plugins."io.containerd.grpc.v1.cri".containerd]

default_runtime_name = "runc"
disable_snapshot_annotations = true
discard_unpacked_layers = false
no_pivot = false
snapshotter = "overlayfs"

[plugins."io.containerd.grpc.v1.cri".containerd.default_runtime]

base_runtime_spec = ""

```
container_annotations = []
pod_annotations = []
privileged_without_host_devices = false
runtime_engine = ""
runtime_root = ""
runtime_type = ""

[plugins."io.containerd.grpc.v1.cri".containerd.default_runtime.options]
```

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes]
```

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc]
  base_runtime_spec = ""
  container_annotations = []
  pod_annotations = []
  privileged_without_host_devices = false
  runtime_engine = ""
  runtime_root = ""
  runtime_type = "io.containerd.runc.v2"
```

```
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.runc.options]
  BinaryName = ""
  CriuImagePath = ""
  CriuPath = ""
  CriuWorkPath = ""
  IoGid = 0
  IoUid = 0
  NoNewKeyring = false
  NoPivotRoot = false
  Root = ""
  ShimCgroup = ""
  SystemdCgroup = true
```

```
[plugins."io.containerd.grpc.v1.cri".containerd.untrusted_workload_runtime]
  base_runtime_spec = ""
  container_annotations = []
  pod_annotations = []
  privileged_without_host_devices = false
```

```
runtime_engine = ""
```

```
runtime_root = ""
```

```
runtime_type = ""
```

```
[plugins."io.containerd.grpc.v1.cri".containerd.untrusted_workload_runtime.options]
```

```
[plugins."io.containerd.grpc.v1.cri".image_decryption]
```

```
key_model = "node"
```

```
[plugins."io.containerd.grpc.v1.cri".registry]
```

```
config_path = ""
```

```
[plugins."io.containerd.grpc.v1.cri".registry.auths]
```

```
[plugins."io.containerd.grpc.v1.cri".registry.configs]
```

```
[plugins."io.containerd.grpc.v1.cri".registry.headers]
```

```
[plugins."io.containerd.grpc.v1.cri".registry.mirrors]
```

```
[plugins."io.containerd.grpc.v1.cri".x509_key_pair_streaming]
```

```
tls_cert_file = ""
```

```
tls_key_file = ""
```

```
[plugins."io.containerd.internal.v1.opt"]
```

```
path = "/opt/containerd"
```

```
[plugins."io.containerd.internal.v1.restart"]
```

```
interval = "10s"
```

```
[plugins."io.containerd.metadata.v1.bolt"]
```

```
content_sharing_policy = "shared"
```

```
[plugins."io.containerd.monitor.v1.cgroups"]
```

```
no_prometheus = false
```

```
[plugins."io.containerd.runtime.v1.linux"]
```

```
no_shim = false
```

```
runtime = "runc"
runtime_root = ""
shim = "containerd-shim"
shim_debug = false
```

```
[plugins."io.containerd.runtime.v2.task"]
  platforms = ["linux/amd64"]
```

```
[plugins."io.containerd.service.v1.diff-service"]
  default = ["walking"]
```

```
[plugins."io.containerd.snapshotter.v1.aufs"]
  root_path = ""
```

```
[plugins."io.containerd.snapshotter.v1.btrfs"]
  root_path = ""
```

```
[plugins."io.containerd.snapshotter.v1.devmapper"]
  async_remove = false
  base_image_size = ""
  pool_name = ""
  root_path = ""
```

```
[plugins."io.containerd.snapshotter.v1.native"]
  root_path = ""
```

```
[plugins."io.containerd.snapshotter.v1.overlayfs"]
  root_path = ""
```

```
[plugins."io.containerd.snapshotter.v1.zfs"]
  root_path = ""
```

```
[proxy_plugins]
```

```
[stream_processors]
```

```
[stream_processors."io.containerd.ocicrypt.decoder.v1.tar"]
  accepts = ["application/vnd.oci.image.layer.v1.tar+encrypted"]
```



```
args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
env = ["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprovider.conf"]
path = "ctd-decoder"
returns = "application/vnd.oci.image.layer.v1.tar"
```

```
[stream_processors."io.containerd.ocicrypt.decoder.v1.tar.gzip"]
accepts = ["application/vnd.oci.image.layer.v1.tar+gzip+encrypted"]
args = ["--decryption-keys-path", "/etc/containerd/ocicrypt/keys"]
env = ["OCICRYPT_KEYPROVIDER_CONFIG=/etc/containerd/ocicrypt/ocicrypt_keyprovider.conf"]
path = "ctd-decoder"
returns = "application/vnd.oci.image.layer.v1.tar+gzip"
```

[timeouts]

```
"io.containerd.timeout.shim.cleanup" = "5s"
"io.containerd.timeout.shim.load" = "5s"
"io.containerd.timeout.shim.shutdown" = "3s"
"io.containerd.timeout.task.state" = "2s"
```

[ttrpc]

```
address = ""
gid = 0
uid = 0
```

Step 8: Restart containerd and Check the Status

```
sudo systemctl restart containerd && systemctl status containerd
```

On Master Node

Step 9: Install kubeadm, kubelet, and kubectl

```
sudo apt-get update
```

```
sudo apt-get install -y apt-transport-https ca-certificates curl gpg
```

```
sudo mkdir -p -m 755 /etc/apt/keyrings
```

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.30/deb/Release.key | sudo gpg --dearmor -o
/etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.30/deb/ /' |
sudo tee /etc/apt/sources.list.d/kubernetes.list
```

```
sudo apt-get update -y
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```

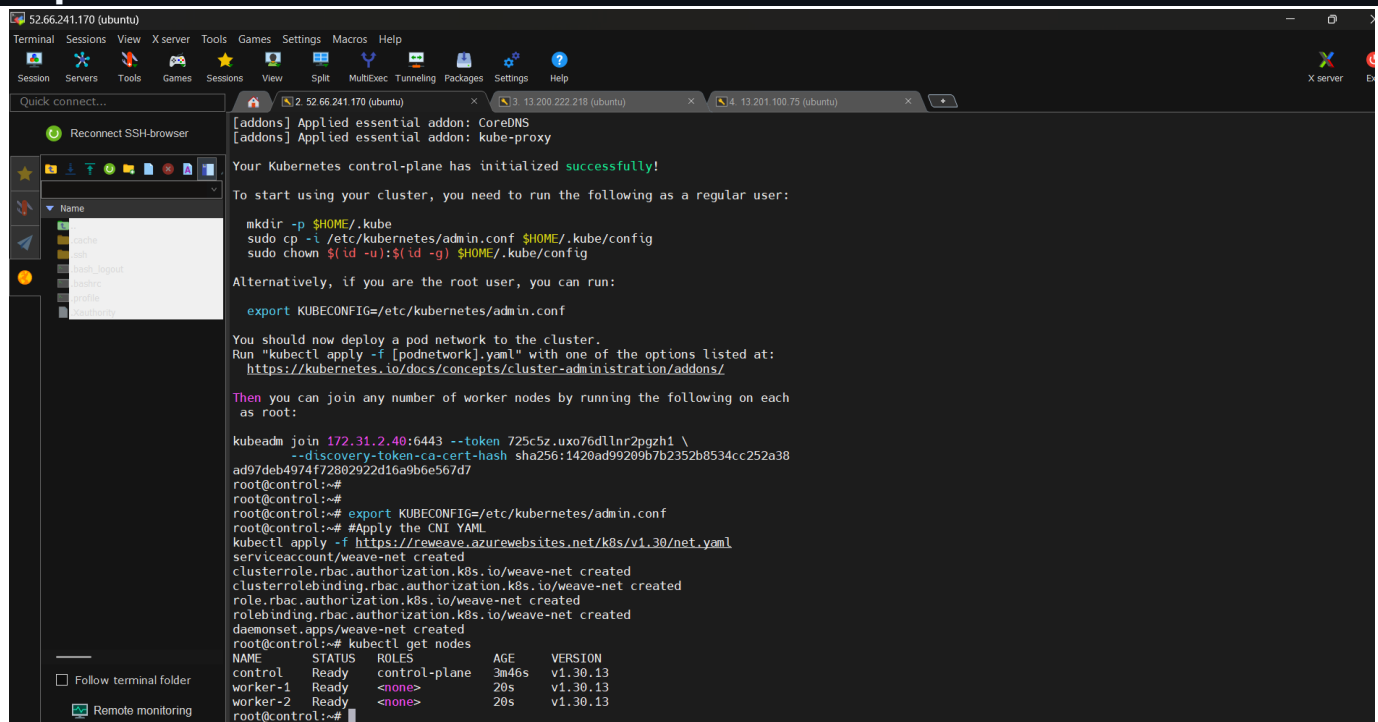
Step 10: Initialize the Cluster and Install CNI

```
sudo kubeadm config images pull
```

```
sudo kubeadm init
```

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

Step 11: Add the worker node to the cluster



```
52.66.241.170 (ubuntu)
Terminal Sessions View X server Tools Games Settings Macros Help
Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help

Quick connect...
Reconnect SSH-browser

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

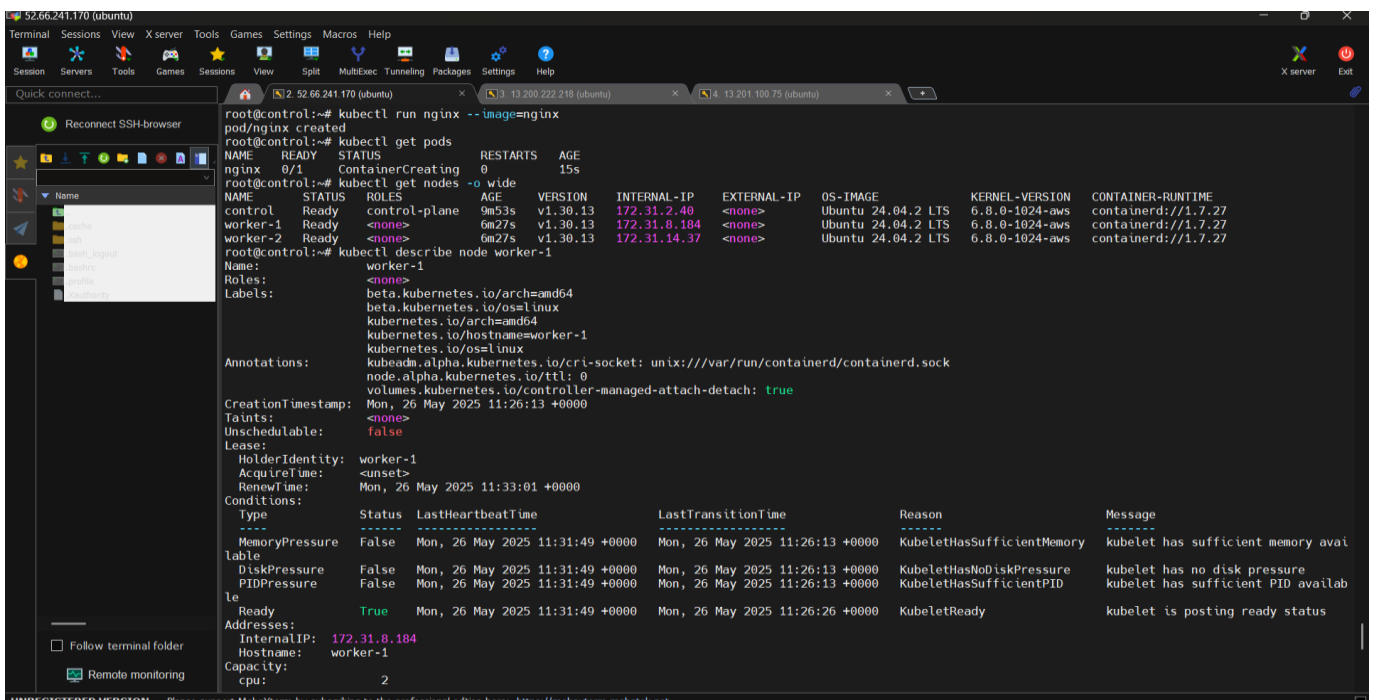
Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each
as root:

kubeadm join 172.31.2.40:6443 --token 725c52.uxo76dllnr2pgzh1 \
--discovery-token-ca-cert-hash sha256:1420ad99209b7b2352b8534cc252a38
ad97deb4974f72802922d16a9b6e567d7
root@control:~#
root@control:~# export KUBECONFIG=/etc/kubernetes/admin.conf
root@control:~# #Apply the CNI YAML
kubectl apply -f https://raw.githubusercontent.com/weaveworks/net/k8s/v1.30/net.yaml
serviceaccount/weave-net created
clusterrole.rbac.authorization.k8s.io/weave-net created
clusterrolebinding.rbac.authorization.k8s.io/weave-net created
role.rbac.authorization.k8s.io/weave-net created
rolebinding.rbac.authorization.k8s.io/weave-net created
daemonset.apps/weave-net created
root@control:~# kubectl get nodes
NAME          STATUS    ROLES    AGE   VERSION
control       Ready     control-plane   3m46s   v1.30.13
worker-1      Ready     <none>        20s     v1.30.13
worker-2      Ready     <none>        20s     v1.30.13
root@control:~#
```



```
52.66.241.170 (ubuntu)
Terminal Sessions View X server Tools Games Settings Macros Help
Session Servers Tools Games Sessions View Split MultiExec Tunneling Packages Settings Help

Quick connect...
Reconnect SSH-browser

root@control:~# kubectl run nginx --image=nginx
pod/nginx created
root@control:~# kubectl get pods
NAME    READY   STATUS    RESTARTS   AGE
nginx   0/1     ContainerCreating   0           15s
root@control:~# kubectl get nodes
NAME          STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE           KERNEL-VERSION   CONTAINER-RUNTIME
control       Ready     control-plane   9m53s   v1.30.13   172.31.2.40    <none>         Ubuntu 24.04.2 LTS   6.8.0-1024-aws   containerd://1.7.27
worker-1      Ready     <none>        6m27s   v1.30.13   172.31.8.184   <none>         Ubuntu 24.04.2 LTS   6.8.0-1024-aws   containerd://1.7.27
worker-2      Ready     <none>        6m27s   v1.30.13   172.31.14.37   <none>         Ubuntu 24.04.2 LTS   6.8.0-1024-aws   containerd://1.7.27
root@control:~# kubectl describe node worker-1
Name:          worker-1
Roles:          <none>
Labels:         beta.kubernetes.io/arch=amd64
                beta.kubernetes.io/os=linux
                kubernetes.io/arch=amd64
                kubernetes.io/hostname=worker-1
                kubernetes.io/os=linux
Annotations:    kubeadm.alpha.kubernetes.io/cri-socket: unix:///var/run/containerd/containerd.sock
                node.alpha.kubernetes.io/ttl: 0
                volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Mon, 26 May 2025 11:26:13 +0000
Taints:         <none>
Unschedulable:  false
Lease:
  HolderIdentity:  worker-1
  AcquireTime:     <unset>
  RenewTime:       Mon, 26 May 2025 11:33:01 +0000
Conditions:
  Type             Status  LastHeartbeatTime             LastTransitionTime             Reason                               Message
  ----             -
  MemoryPressure    False   Mon, 26 May 2025 11:31:49 +0000 Mon, 26 May 2025 11:26:13 +0000 KubeletHasSufficientMemory         kubelet has sufficient memory avail
  DiskPressure      False   Mon, 26 May 2025 11:31:49 +0000 Mon, 26 May 2025 11:26:13 +0000 KubeletHasNoDiskPressure           kubelet has no disk pressure
  PIDPressure       False   Mon, 26 May 2025 11:31:49 +0000 Mon, 26 May 2025 11:26:13 +0000 KubeletHasSufficientPID            kubelet has sufficient PID availab
  Ready             True    Mon, 26 May 2025 11:31:49 +0000 Mon, 26 May 2025 11:26:26 +0000 KubeletReady                       kubelet is posting ready status
Addresses:
  InternalIP:  172.31.8.184
  Hostname:     worker-1
Capacity:
  cpu: 2
```